

## Memory Management System

### Virtual Memory

A technique which allows the execution of processes that may not be completely in memory

## Memory Management System

### Fact:

In many cases, for the execution of a program, one does not need to have the entire program in the main memory

### Examples:

- Codes for those options in an interface not chosen by user.
- Part of the process that handles unusual error conditions
- A good portion of an unused large declared arrays

## Memory Management System

### Virtual Memory Adoption

- Overlay
- Dynamic Loading

## Memory Management System

### Overlay Concept

1. Keep those pages of logical space that are essential to execution of a process all the time in memory and
2. Load other pages of the logical space when they are needed and load a needed page into space that was previously occupied by a page that is no longer needed.

## Memory Management System

### Basic Overlay Implementation


The total # of pages that make a process,  $n = 7$ ;

The total number of frames given to the process,  $k = 4$ ;

The total number of pages that absolutely must be in,  $m = 3$ ;



Logical Space (LS) at  $t_0$



## Memory Management System


### Basic Overlay Implementation

At time  $t_1$ :


Say, P2 is needed

At time  $t_2$ :


Say, P4 is needed




Logical Space (LS) at  $t_0$



LS at  $t_1$



LS at  $t_2$



## Memory Management System

### Dynamic Loading

The same as the Overlay technique with one difference:

Only the main module is loaded and the other modules are loaded upon request by CPU.

Page Table

Physical Memory

## Memory Management System

The most common virtual memory system is  
Demand Paging System

Page Table

Physical Memory

## Memory Management System

### Demand Paging System

Paging + Swapping

∴ We need :

Page Table + Backing Store


Swapping is done by a lazy swapper (a.k.a. Pager). A swapper is called lazy because it swaps in/swaps out only needed pages. (A regular Memory swapper may swap needed pages + extra pages)

## Demand Paging System

### Assumptions:

- A process  $P$  is given and it has only 7 pages ( $P_0, P_6$ ) that are initially located on the backing store.

Physical Memory



Backing Store

## Demand Paging System

### Assumptions:

- 7 pages are designated as the logical memory space. Therefore, Page Table has 7 entries

F#	Validity Bit

Page Table

Physical Memory



Backing Store

## Demand Paging System


### Assumptions:

- Only three pages of  $P_0$ ,  $P_2$ , and  $P_4$  of the logical memory space, that are absolutely needed are loaded in the physical memory.

F#	Validity Bit
2	
5	
9	

Page Table

Physical Memory



Backing Store

## Memory Management System

Frames: bit Protection Validity

0	F2	00	1
1	F9	01	1
2	F8	00	1
3	F12	11	1
4			0
5			0
6			0
7			0

Validity bit says: whether the entry is empty (validity = 0) or not

We add an extra meaning to the validity bit for use in DEMEND PAGING

Validity bit says:

1. whether the entry is valid or not and the associated page is legal (i.e. page# verification that it belongs to the logical space is already completed.)
- 2.


## Demand Paging System

### Assumptions:


- Only three pages of P0, P2, and P4 of the logical memory space, that are absolutely needed are loaded in the physical memory.

F#	Validity Bit
2	1
	0
5	1
	0
9	1
	0
	0

Page Table



Physical Memory



Backing Store

## Demand Paging System


F# Validity Bit

2	1
	0
5	1
	0
9	1
	0
	0

Page Table



Physical Memory



Backing store

## Demand Paging System

### Error is trapped

1. The status of the process that caused the page fault is saved.
2. Since valid bit = 0 an Internal Table, kept with the PCB--Process Control Block, determines whether such page is legal or not.
3. If illegal

Then: terminate the process.

Else:


- OS gets the demanded page from the backing store
- A free frame is found by OS
- Page is loaded in the physical memory
- Page table and Internal Table is updated
- Process is re-started and continues.

## Demand Paging System


F# Validity Bit

2	1
7	1
5	1
	0
9	1
	0
	0

Page Table



Physical Memory



Backing store

## Memory Management System

### Problem:

(Sometimes) it is difficult to start execution of a process exactly from the point that caused the page fault.

Page Table

Physical Memory

## Memory Management System

### Example 1:

Page fault happens In the middle of a 3-address Instruction, like  $C=A+B$

fetch A; Fetch B; Add A,B; Store in C;

Page fault happens locating C.

After Page fault the addition must be started from the beginning

Page Table

Physical Memory

## Memory Management System

### Example 2:

Page fault happens In the middle of a Move instruction in PDP-11:

MOVE (R2)\*, -(R3)

This instruction means:

After R2 is used as a pointer increment R2 content by 2.

Decrement R3 content by 2 before use it as a pointer.

Page fault happens because R3 now pointing to an address in a new page

After the page fault the Move must be started from the beginning. The contents of both R2 and R3 have been already destroyed

Page Table

Physical Memory

## Memory Management System

### Solutions:

1. Save the content of registers in another set of registers for use in case a page fault happens
2. Examine all the addresses involved in an instruction before executing the instruction (How could it be done? Using micro code). If page faults are needed, do it before actually execute the instructions.

Page Table

Physical Memory

## Memory Management System

### Performance of Demand Paging

Demand paging is pure overhead and it degrades the performance of the system.

### Question:

How many page fault should we have, if the goal is to have <10% degradation in system performance?

Page Table

Physical Memory

## Memory Management System

### Performance of Demand Paging

Answer: Let  $p$  be the probability of page fault  
 $m_a$  be the access time for the page table  
 $PF$  be the time completing a page fault.

Effective Access time,  $EA = p(PF) + (1-p)m_a$

For  $PF = 10ms$  and  $m_a = 1 \mu s$ ,  $EA = 10000p + 1 - p = 1 + 9999p \mu s$   
 where 9999p is the degradation factor

10% degradation means  $9999p < 0.10$ ;  $p < 0.1/9999$ ;  $p \approx 0.00001$ ;

It means: to have <10% degradation, the probability of having a page fault must be 1 out of 100,000 demand paging.

Physical Memory

## Memory Management System

Degree of multiprogramming (DM) is influenced by the number of frames, X, allocated to each process

$$DM = \lfloor \text{total frames in physical memory} / X \rfloor$$

Problem #1: What if the process needs more than X frames?

Problem #2: How did we arrive at X? (allocation Problem)

Page Table

Physical Memory

## Memory Management System

Answer to Problem #1: Use of Page Replacement Technique

Page Replacement Technique:

- 1- Find a frame (victim frame)
- 2- Write the content of the frame to the backing store
- 3- Update the page table and internal table
- 4- The freed frame now can be used as part of the page fault handling.

Note: Use of a dirty bit eliminates the second step.

Page Table

Physical Memory

## Memory Management System

Page Replacement Algorithms

Given:

Reference String: A string of memory references  
(Created artificially or by tracing a system)

Number of Frames given to the process  
(Higher the number of frames, lower the number of page faults)

Page Table

Physical Memory

## Memory Management System

Let say a process is made up of 50 pages. However, only 4 frames are dedicated to the process.

The page table has 50 entries. At any given time maximum of four entries of the page table are full. These entries are related to the four pages that currently occupy the four frames (in RAM)

During the discussion of the algorithms we see something like the following picture: This picture shows the four frames of RAM that in reality may not be contiguous but we show them in contiguous form for ease of drawing. The picture says that the 4 frames have been occupied by page 7, page 3, page 2, and page 4.

This picture does not have anything to do with the page table except that the entries of 2, 3, 4, and 7 of the page table have frame numbers and the rest of the entries are empty

## Memory Management System

Page Replacement Algorithms

1. FIFO (Replace the page that is at the head of queue)
2. Optimal Replacement (Replace the page that will not be used for the longest period of time)
3. LRU (Least Recently Used)
4. LRU Approximation
5. Second Chance Replacement
6. LFU (Least Frequently Used)
7. MFU (Most Frequently Used)
8. Page Classes
9. Ad Hoc Algorithms

Page Table

Physical Memory

## Memory Management System

Page Replacement Algorithms

FIFO (replace the page that is at the head of queue)

Reference String: 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0, 1

7	7	7	0	1	2	3	0	4	2	3	0	1	2	7
0	0	1	2	3	0	4	2	3	0	1	2	7	0	
1	2	3	0	4	2	3	0	1	2	7	0	1		

15 page Replacement

- Belady's Anomaly: Reference String: 1, 2, 3, 4, 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

Using 4 frames

1	1	1	1	2	3	4	5	1	2
2	2	2	3	4	5	1	2	3	10
3	3	4	5	1	2	3	4	P.R.	
4	5	1	2	3	4	5			

Using 3 frames

1	1	1	2	3	4	1	2	5	3
2	2	3	4	1	2	5	3	9	
3	4	1	2	5	3	4	P.R.		

## Memory Management System

Optimal Replacement (replace the page that will not be used for the

Reference String: 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0, 1

7	7	7	2	2	2	2	7
0	0	0	0	4	0	0	0
1	1	3	3	1	1		

9 page Replacement

LRU (Least Recently Used)

7	7	7	2	2	4	4	0	0	0
0	0	0	0	0	0	3	3	3	3
1	1	3	3	2	2	2	2	2	2

12 page Replacement

## Memory Management System

### LRU Implementation

Using time-of-use (Clock)

Using a double-linked list

## Memory Management System

### LRU Implementation

Using time-of-use (Clock)

- A logical clock (a Counter) is used and set to zero Counter = 0
- Clock will be incremented by one when a page reference (in Reference string) is made.
- The new clock value is recorded for the referenced page.
- Always the page with smallest value of the counter is the least recent used page.

## Memory Management System

### LRU Implementation

- Using time-of-use

Example: Clock = 0 Reference String: 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0, 1  
 Referencing 7: Clock = 1, Referencing 0: Clock = 2, Referencing 1: Clock = 3,  
 Referencing 2: Clock = 4 and page 7 with smallest clock is replaced.  
 Referencing 0: Clock = 5 and no replacement. Referencing 3: Clock = 6 and page 1 is replaced. Referencing 0: Clock = 7 and no replacement (clock is updated for P=0.)  
 Referencing 4: Clock = 8 and page 2 with smallest clock value is replaced.

P#	C	P#	C	P#	C	P#	C	P#	C	P#	C	P#	C	P#	C	P#	C	P#	C
7	1	7	1	7	1	2	4	2	4	2	4	2	4	4	8				
		0	2	0	2	0	2	0	5	0	5	0	7	0	7				
				1	3	1	3	1	3	3	6	3	6	3	6				

## Memory Management System

### LRU Implementation

- Use of double-linked list.

There is a head and tail pointer. The most recent used page is the head of the list and time that a page is referenced it is moved to the head of list. Therefore, tail pointer always point to the LRU page.

Page Table

Physical Memory

## Memory Management System

### LRU Approximation

- 1- Use of reference bit.

Each page has its own reference bit.

The bit is initialized with zero when the page is coming to the physical Memory.

The bit is set to one if it is referenced again.

The page with reference bit = 0 is the least recently used page.

Reference string: 7, 1, 7, 2 and we have 3 frames

P#	P#	R	P#	R	P#	R	P#	R	P#	R	P#	R
0	7	0	7	0	7	0	7	1	7	1		
1		0	0	0	0	0	0	2	0			
2				1	0	1	0	1	0			

(Either of the two pages 1 or 2 may be replaced)

## Memory Management System

Reference bit neither says anything about the history of the page nor the order of page use.

(That is the reason for calling it LRU Approximation)

Page Table

## Memory Management System

### LRU Approximation

1- Use of reference byte. Each page has its own reference byte. The byte is initialized with 11111111 when the page is coming to the main memory.

## Memory Management System

### LRU Approximation

1- Use of reference byte. Each page has its own reference byte. The byte is initialized with 11111111 when the page is coming to the main memory.

$t = 0$

Frame #	Reference byte
0 2	11111111
1 5	11111111
2 12	11111111

Page Table

Physical Memory

## Memory Management System

### LRU Approximation

1- Use of reference byte. Each page has its own reference byte. The byte is initialized with 11111111 when the page is coming to the main memory.

$t = 0$

Frame #	Reference byte
0 2	11111111
1 5	11111111
2 12	11111111

After each fixed period of time (say 100  $\mu$ s) every reference byte is shift one bit to the right.

If the page is not referenced since the last shift

Then, The empty bit is set to zero;  
Otherwise, it is set to 1.

Physical Memory

## Memory Management System

### LRU Approximation

1- Use of reference byte. Each page has its own reference byte. The byte is initialized with 11111111 when the page is coming to the main memory.

After the first fixed period of time (say 100  $\mu$ s), only page 1 has been referenced:

$t = 0$

Frame #	Reference byte
0 2	11111111
1 5	11111111
2 12	11111111

Page Table

Frame #	Reference byte
0 2	01111111
1 5	11111111
2 12	01111111

Physical Memory

## Memory Management System

### LRU Approximation

1- Use of reference byte. Each page has its own reference byte. The byte is initialized with 11111111 when the page is coming to the main memory.

$t = 0$

Frame #	Reference byte
0 2	11111111
1 5	11111111
2 12	11111111

$t = t_0$

Frame #	Reference byte
0 2	00100111
1 5	10000110
2 12	00100101

The page with the smallest unsigned integer number in its reference byte is the LRU page.


Physical Memory

## Memory Management System

### Second Chance Replacement

Pages in the page table make a circular queue and a pointer points to the next victim in the queue. To each page a reference bit is associated that initially is set to zero. Upon the first reference to a page after the initial loading the reference bit is set to one.

Pointed page is the candidate for replacement.



If the reference bit of the pointed page is 1 it is set to 0, pointer is incremented, and page will not be replaced.

The process continues until the first page with the reference bit = 0 is found.

That page will be replaced.