

Memory Management System

Virtual Memory

A technique which allows the execution of processes that may not be completely in memory

Memory Management System

Fact:

In many cases, for the execution of a program, one does not need to have the entire program in the main memory

Examples:

- Codes for those options in an interface not chosen by user.
- Part of the process that handles unusual error conditions
- A good portion of an unused large declared arrays

Memory Management System

Virtual Memory Adoption

- Overlay
- Dynamic Loading

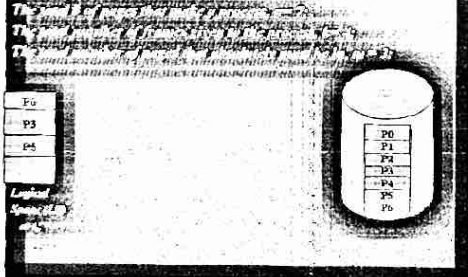
Memory Management System

Overlay Concept

1. Keep those pages of logical space that are essential to execution of a process all the time in memory and
2. Load other pages of the logical space when they are needed and load a needed page into space that was previously occupied by a page that is no longer needed.

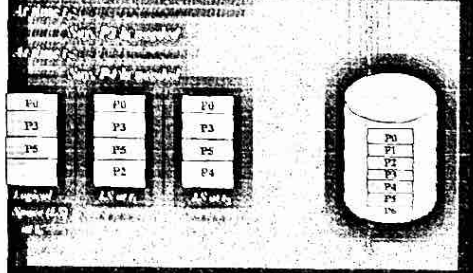
Memory Management System

Basic Overlay Implementation



Memory Management System

Basic Overlay Implementation



Memory Management System

Dynamic Loading

The same as the Overlay technique with one difference:
Only the main module is loaded and the other modules
are loaded upon request by CPU.

Page Table

Physical Memory

Memory Management System

The most common virtual memory system is
Demand Paging System

Page Table

Physical Memory

Memory Management System

Demand Paging System

Paging + Swapping

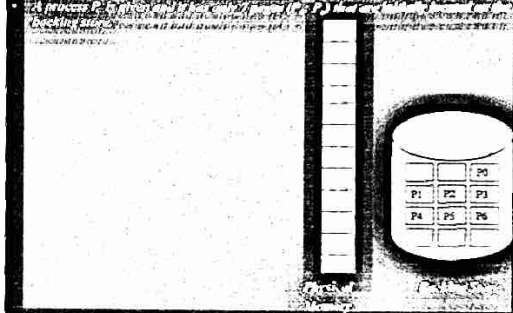
∴ We need :

Page + Backing
Table Store

Swapping is done by a lazy swapper (a.k.a. Pager). A swapper is called
lazy because it swaps in/swaps out only needed pages. (A regular Memory
swapper may swap needed pages + extra pages)

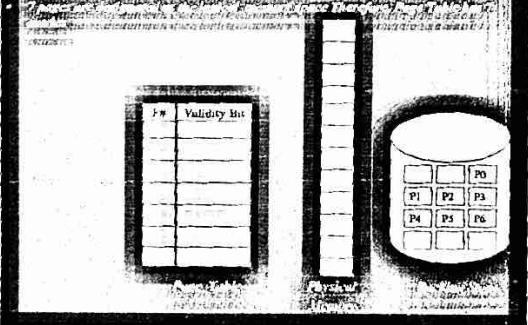
Demand Paging System

Assumptions:



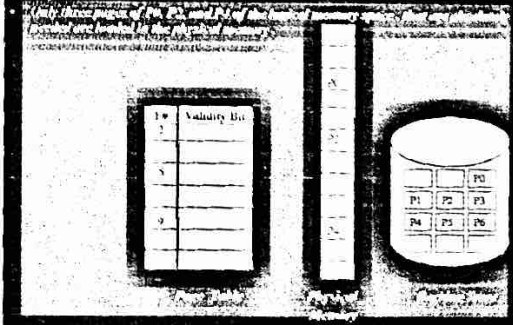
Demand Paging System

Assumptions:



Demand Paging System

Assumptions:



Memory Management System

Process	Page	Valid Bit
0	F2	00
1	F9	01
2	F8	00
3	F12	11
4		0
5		0
6		0
7		0

Validity bit says: whether the entry is empty (validity = 0) or not

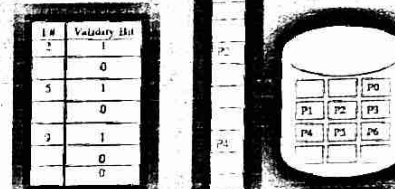
We add an extra meaning to the validity bit for use in DEMEND PAGING

- Validity bit says:
- whether the entry is valid or not and the associated page is legal (i.e. page# verification that it belongs to the logical space is already completed)

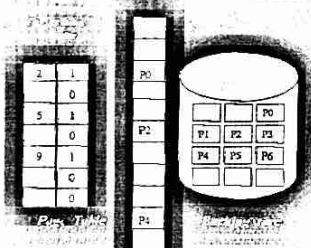
Demand Paging System

Assumptions:

- Each process has its own PCB (Process Control Block)
- Each process has its own internal table
- Each process has its own physical memory



Demand Paging System

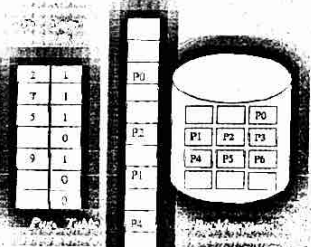


Demand Paging System

Error is trapped

- The status of the process that caused the page fault is saved
 - Since valid bit = 0 in internal table, entry with the PCB - Process Control Block determines whether such page is legal or not
 - If illegal, the process is terminated
- Otherwise:
- OS still has to find the physical page
 - OS will find the physical page
 - Page is loaded in the physical memory
 - Change table and internal table is
 - Process is re-started and continues

Demand Paging System



Memory Management System

Problem:

(Sometimes) it is difficult to start execution of a process exactly from the point that caused the page fault.

Page Table

Physical Memory

Memory Management System

Example 1:

Page fault happens in the middle of a 3-address instruction, like $C=A+B$

Fetch A; Fetch B; Add A,B; Store in C;

Page fault happens locating C.

After Page fault the addition must be started from the beginning

Page Table

Physical Memory

Memory Management System

Example 2:

Page fault happens in the middle of a Move instruction in PDP-11:

MOVE (R2), -(R3)

This instruction means:

After R2 is used as a pointer increment R2 content by 2.

Decrement R3 content by 2 before use it as a pointer.

Page fault happens because R3 now pointing to an address in a new page

After the page fault the Move must be started from the beginning. The contents of both R2 and R3 have been already destroyed.

Page Table

Physical Memory

Memory Management System

Solutions:

1. Save the content of registers in another set of registers for use in case a page fault happens
2. Examine all the addresses involved in an instruction before executing the instruction (How could it be done? Using micro code). If page faults are needed, do it before actually execute the instructions.

Page Table

Physical Memory

Memory Management System

Performance of Demand Paging

Demand paging is pure overhead and it degrades the performance of the system.

Question:

How many page fault should we have, if the goal is to have <10% degradation in system performance?

Page Table

Physical Memory

Memory Management System

Performance of Demand Paging

Answer: Let p be the probability of page fault
 m_a be the access time for the page table
 PF be the time completing a page fault

Effective Access time, $EA = p(PF) + (1-p)m_a$

For $PF = 10ms$ and $m_a = 1 \mu s$, $EA = 10000p + 1 - p = 1 + 9999p \mu s$
 where 9999p is the degradation factor

10% degradation means $9999p < 0.10$; $p < 0.1/9999$; $p \approx 0.00001$;
 It means: to have <10% degradation, the probability of having a page fault must be 1 out of 100,000 demand paging.

Physical Memory

Memory Management System

Degree of multiprogramming (DM) is influenced by the number of frames, N , allocated to each process

$DM = \frac{1}{N} \times (\text{total frames in physical memory})$

- Problem #1: What if the process needs more than N frames?
 Problem #2: How did we arrive at N ? (allocation Problem)

Page Table

Physical Memory

Memory Management System

Answer to Problem #1: Use of Page Replacement Technique

Page Replacement Technique:

- 1- Find a frame (victim frame)
- 2- Write the content of the frame to the backing store
- 3- Update the page table and internal table
- 4- The freed frame now can be used as part of the page fault handling.

Note: Use of a dirty bit eliminates the second step.

Page Table

Physical Memory

Memory Management System

Page Replacement Algorithms

Given:

Reference String: A string of memory references
(Created artificially or by tracing a system)

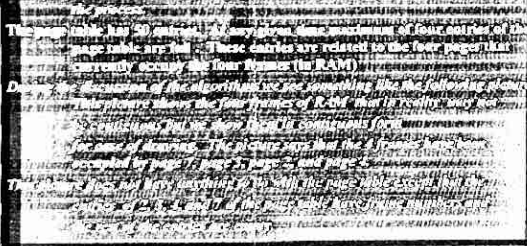
Number of Frames given to the process
(Higher the number of frames, lower the number of page faults)

Page Table

Physical Memory

Memory Management System

Let say a process is made up of 30 pages. However, only 4 frames are dedicated to



Memory Management System

Page Replacement Algorithms

1. FIFO (Replace the page that is at the head of queue)
2. Optimal Replacement (Replace the page that will not be used for the longest period of time)
3. LRU (Least Recently Used)
4. LRU Approximation
5. Second Chance Replacement
6. LFU (Least Frequently Used)
7. MFU (Most Frequently Used)
8. Page Classes
9. Ad Hoc Algorithms

Page Table

Physical Memory

Memory Management System

Page Replacement Algorithms

FIFO (replace the page that is at the head of queue)

Reference String: 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0, 1

7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7	0	1
0	0	1	2	3	0	4	2	3	0	1	2	7	0						
1	2	3	0	4	2	3	0	1	2	7	0	1							

15 page Replacement

- Belady's Anomaly: Reference String: 1, 2, 3, 4, 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

Using 4 frames

1	2	3	4	5	1	2	3	4	5	1	2	3	4	5
2	2	3	4	5	1	2	3	10						
3	3	4	5	1	2	3	4	P.R.						
4	5	1	2	3	4	5								

Using 3 frames

1	2	3	4	5	1	2	3	4	5	1	2	3	4	5
2	2	3	4	1	2	5	3	9						
3	4	1	2	5	3	4	P.R.							

Memory Management System

Optimal Replacement (replace the page that will not be used for the

Reference String: 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0, 1

7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7	0	1
0	0	0	0	4	0	0													
1	1	3	3	1	1														

9 page Replacement

LRU (Least Recently Used)

7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7	0	1
0	0	0	0	0	0	3	3	3	3										
1	1	3	3	2	2	2	2	2	2										

12 page Replacement

Memory Management System

LRU Implementation

Using time-of-use (Clock)
Using a double-linked list

Memory Management System

LRU Implementation

Using time-of-use (Clock)

- A logical clock (a Counter) is used and set to zero Counter = 0
- Clock will be incremented by one when a page reference (in Reference string) is made.
- The new clock value is recorded for the referenced page.
- Always the page with smallest value of the counter is the least recent used page.

Memory Management System

LRU Implementation

- Using time-of-use

Example: Clock = 0 Reference String: 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0, 1
Referencing 7: Clock = 1, Referencing 0: Clock = 2, Referencing 1: Clock = 3,
Referencing 2: Clock = 4 and page 7 with smallest clock is replaced.
Referencing 0: Clock = 5 and no replacement. Referencing 3: Clock = 6 and page 1 is replaced. Referencing 0: Clock = 7 and no replacement (clock is updated for P=0).
Referencing 4: Clock = 8 and page 2 with smallest clock value is replaced.

P#	C	P#	C	P#	C	P#	C	P#	C	P#	C	P#	C	P#	C	P#	C	P#	C
7	1	7	1	7	1	2	4	2	4	2	4	2	4	4	4	8			
		0	2	0	2	0	2	0	5	0	5	0	7	0	7				
				1	3	1	3	1	3	3	6	3	6	3	6				

Memory Management System

LRU Implementation

- Use of double-linked list.

There is a head and tail pointer. The most recent used page is the head of the list and time that a page is referenced it is moved to the head of list. Therefore, tail pointer always point to the LRU page.

Page Table

Physical Memory

Memory Management System

LRU Approximation

- Use of reference bit.

Each page has its own reference bit.

The bit is initialized with zero when the page is coming to the physical Memory.

The bit is set to one if it is referenced again.

The page with reference bit = 0 is the least recently used page.

Reference string: 7, 1, 7, 2 and we have 3 frames

P#	P#	R	P#	R	P#	R	P#	R	P#	R	P#	R	P#	R	P#	R	P#	R	P#	R
0	7	0	7	0	7	0	7	1	7	1										
1		0	0	0	0	0	0	0	2	0										
2				1	0	1	0	1	0											

(Either of the two pages 1 or 2 may be replaced)

Memory Management System

Reference bit neither says anything about the history of the page nor the order of page use.

(That is the reason for calling it LRU Approximation)

Page Table

Memory Management System

LRU Approximation

1- Use of reference byte. Each page has its own reference byte. The byte is initialized with 11111111 when the page is coming to the main memory.

Memory Management System

LRU Approximation

1- Use of reference byte. Each page has its own reference byte. The byte is initialized with 11111111 when the page is coming to the main memory.

$t = 0$

Frame #	Reference byte
0 2	11111111
1 5	11111111
2 12	11111111

Page Table

Physical Memory

Memory Management System

LRU Approximation

1- Use of reference byte. Each page has its own reference byte. The byte is initialized with 11111111 when the page is coming to the main memory.

$t = 0$

Frame #	Reference byte
0 2	11111111
1 5	11111111
2 12	11111111

After each fixed period of time (say 100 μ s) every reference byte is shift one bit to the right.

If the page is not referenced since the last shift

Then, The empty bit is set to zero;

Otherwise, it is set to 1.

Physical Memory

Memory Management System

LRU Approximation

1- Use of reference byte. Each page has its own reference byte. The byte is initialized with 11111111 when the page is coming to the main memory.

After the first fixed period of time (say 100 μ s), only page 1 has been referenced:

$t = 0$

Frame #	Reference byte
0 2	11111111
1 5	11111111
2 12	11111111

$t = t_1$

Frame #	Reference byte
0 2	01111111
1 5	11111111
2 12	01111111

Page Table

Physical Memory

Memory Management System

LRU Approximation

1- Use of reference byte. Each page has its own reference byte. The byte is initialized with 11111111 when the page is coming to the main memory.

$t = 0$

Frame #	Reference byte
0 2	11111111
1 5	11111111
2 12	11111111

$t = t_n$

Frame #	Reference byte
0 2	00100111
1 5	10000110
2 12	00100101

The page with the smallest unsigned integer number in its reference byte is the LRU page.

Physical Memory

Memory Management System

Second Chance Replacement

Pages in the page table make a circular queue and a pointer points to the next victim in the queue. To each page a reference bit is associated that initially is set to zero. Upon the first reference to a page after the initial loading the reference bit is set to one.

Pointed page is the candidate for replacement

If the reference bit of the pointed page is 1 it is set to 0, pointer is incremented, and page will not be replaced.

The process continues until the first page with the reference bit = 0 is found.

That page will be replaced.

